

B Trees

Prof. Zheng Zhang

Harbin Institute of Technology, Shenzhen



OUTLINE

- *Introduction*
- Definition
- Basic Operations
- Applications



INTRODUCTION

- Motivation
- Overview
- Original Publication
- B-tree ...



INTRODUCTION--MOTIVATION

- When data is too large to fit in main memory, then the number of disk accesses becomes important.
- Disk access is unbelievably expensive compared to a typical computer instruction (mechanical limitations).
- One disk access is worth about 200,000 instructions.
- The number of disk accesses will dominate the running time.



INTRODUCTION--MOTIVATION

- Secondary memory (disk) is divided into equal-sized blocks (typical sizes are 512, 2048, 4096 or 8192 bytes).
- Basic I/O operation transfers the contents of one disk block to/from main memory.
- Goal is to devise a multiway search tree that will **minimize file accesses** (by exploiting disk block read).



OVERVIEW

- Have many children, from a handful to thousands.
- The "**branching factor**" of a B-tree can be **quite large**, although it is usually determined by characteristics of the disk unit used.
- B-trees are *similar* to red-black trees in that every n -node B-tree has height $O(\lg n)$.
- Although the height of a B-tree can be considerably less than that of a red-black tree because its branching factor can be much larger.




The diagram illustrates a recursive tree structure for a B-tree. The root node is labeled $T.root$ and contains 1000 keys. It has three children, each also containing 1000 keys. The middle child has three children of its own. Each node is represented by a rectangle with a semi-circle at the bottom, and arrows indicate the pointers to the children. The diagram shows the growth of the tree as more nodes are added, with the total number of nodes and keys increasing significantly.

Level	Number of Nodes	Total Number of Keys
Root	1	1000
Level 1	3	3000
Level 2	9	9000
Level 3	27	27000
Level 4	81	81000
Level 5	243	243000
Level 6	729	729000
Level 7	2187	2187000
Level 8	6561	6561000
Level 9	19683	19683000
Level 10	59049	59049000
Level 11	177147	177147000
Level 12	531441	531441000
Level 13	1594323	1594323000
Level 14	4782969	4782969000
Level 15	14348907	14348907000
Level 16	43046721	43046721000
Level 17	129140163	129140163000
Level 18	387420489	387420489000
Level 19	1162261467	1162261467000
Level 20	3486784401	3486784401000
Level 21	10460353203	10460353203000
Level 22	31381059609	31381059609000
Level 23	94143178827	94143178827000
Level 24	282429536481	282429536481000
Level 25	847288609443	847288609443000
Level 26	2541865828329	2541865828329000
Level 27	7625597484987	7625597484987000
Level 28	22876792454961	22876792454961000
Level 29	68630377364883	68630377364883000
Level 30	205891132094649	205891132094649000
Level 31	617673396283947	617673396283947000
Level 32	1853020188851841	1853020188851841000
Level 33	5559060566555523	5559060566555523000
Level 34	16677181699666569	16677181699666569000
Level 35	50031545098999707	50031545098999707000
Level 36	150094635296999121	150094635296999121000
Level 37	450283905890997363	450283905890997363000
Level 38	1350851717672992089	1350851717672992089000
Level 39	4052555153018976267	4052555153018976267000
Level 40	12157665459056928801	12157665459056928801000
Level 41	36472996377170786403	36472996377170786403000
Level 42	109418989131512359209	109418989131512359209000
Level 43	328256967394537077627	328256967394537077627000
Level 44	984770902183611232881	984770902183611232881000
Level 45	2954312706550833698643	2954312706550833698643000
Level 46	8862938119652501095929	8862938119652501095929000
Level 47	26588814358957503287787	26588814358957503287787000
Level 48	79766443076872509863361	79766443076872509863361000
Level 49	239299329230617529580083	239299329230617529580083000
Level 50	717897987691852588740249	717897987691852588740249000
Level 51	2153693963075557766220747	2153693963075557766220747000
Level 52	6461081889226673308662241	6461081889226673308662241000
Level 53	19383245667680019925986723	19383245667680019925986723000
Level 54	58149737003040059777960169	58149737003040059777960169000
Level 55	174449211009120179333880507	174449211009120179333880507000
Level 56	523347633027360537991641521	523347633027360537991641521000
Level 57	1570042899082081613974924563	1570042899082081613974924563000
Level 58	4710128697246244841924773689	4710128697246244841924773689000
Level 59	14130386091738734525774321067	14130386091738734525774321067000
Level 60	42391158275216203577322963201	42391158275216203577322963201000
Level 61	127173474825648610731968889603	12717347482564861073196888960



EXAMPLE


- For a large B-tree stored on a disk, branching factors between 50 and 2000 are *often* used, depending on the size of a key relative to the size of a page.
 - A large branching factor *dramatically reduces* both the *height* of the tree and the *number* of disk accesses required to **find any key**.
 - This B-tree with a *branching factor of 1001* and height 2 that can store over one **billion** keys.
 - Nevertheless, since the *root* node can be kept permanently in main memory, *only two disk accesses* at most are required to find any key in this tree!
- 

OVERVIEW--STRUCTURES

- If an **internal** B-tree node x contains $n[x]$ keys, then x has $n[x]+1$ children.
- The keys in node x are used as dividing points separating the range of keys handled by x into $n[x]+1$ subranges, each handled by one child of x .
- When searching for a key in a B-tree, we make an $(n[x]+1)$ -way decision based on comparisons with the $n[x]$ keys stored at node x .
- The *structure of leaf nodes* differs from that of internal nodes; we will examine these differences later.



EXAMPLE



OVERVIEW--HOW TO WORK

- In a typical B-tree application, the amount of data handled is so large that all the data do not fit into main memory at once.
- The B-tree algorithms copy selected pages from disk into main memory as **needed** and write back onto disk the pages that have **changed**.
- B-tree algorithms are designed so that only a constant number of pages are in main memory at any time.
- Thus, the size of main memory does not limit the size of B-trees that can be handled.



OVERVIEW--THE ORIGINAL PUBLICATION

- Rudolf Bayer, Edward M. McCreight
- Organization and Maintenance of Large Ordered Indices
- 1972



OVERVIEW--B-TREE...

- **Also known as** balanced multiway tree
- **Generalization** (I am a kind of ...)
balanced tree, search tree.
- **Specialization** (... is a kind of me.)
2-3-4 tree, B-tree, 2-3 tree . . .*



OVERVIEW--B-TREE...

- The B-tree's creators, Rudolf Bayer and Ed McCreight, have not explained what, if anything, the **B stands for**.
- The most common belief is that **B stands for *balanced***, as all the leaf nodes are at the same level in the tree.
- B may also stand for *Bayer*, or for *Boeing*, because they were working for Boeing Scientific Research Labs at the time.



OUTLINE


- Introduction
- *Definition*
- Basic Operations
- Applications



DEFINITION

A *B-tree* T is a rooted tree having the following properties:


- **1. Every node x has the following fields**
 - $n[x]$, the number of keys currently stored in node x .
 - The $n[x]$ keys themselves stored in **nondecreasing** order, so that $\text{key}_1[x] \leq \text{key}_2[x] \leq \dots \leq \text{key}_{n[x]}[x]$.
 - $\text{Leaf}[x]$, a boolean value that is TRUE if x is a leaf and FALSE if x is an internal node.



DEFINITION

A *B-tree* T is a rooted tree having the following properties:


- **2.** Each internal node x also contains $n[x]+1$ pointers $c_1[x], c_2[x], \dots, c_{n[x]+1}[x]$ to its children. Leaf nodes have no children, so their c_i fields are undefined.
- **3.** The keys $key_i[x]$ separate the ranges of keys stored in each subtree: if k_i is any key stored in the subtree with the root $c_i[x]$, then
$$k_1 \leq key_1[x] \leq k_2 \leq key_2[x] \leq \dots \leq key_{n[x]}[x] \leq k_{n[x]+1}$$



DEFINITION

A **B-tree** T is a rooted tree having the following properties:

- **4.** All leaves have the same depth, which is the tree's height h .
- **5.** There are lower and upper bounds on the number of keys that a node can contain. These bounds can be expressed in terms of a fixed integer $t \geq 2$ called a *minimum degree* of the B-tree:



DEFINITION

A *B-tree* T is a rooted tree having the following properties:

- **5.** There are lower and upper bounds on the number of keys that a node can contain. These bounds: a minimum degree of the B-tree:
 - Every node other than the root must have at least $t-1$ keys.
Every internal node other than the root thus has at least t *children*. If the tree is **nonempty**, **the root** must have at least one key.
 - Every node can contain at most $2t-1$ keys, therefore, an internal node can have at most $2t$ *children*. we say that a node is *full* if it contains exactly $2t-1$ keys.



DEFINITION--B-TREE HEIGHT


Depth Nodes

0 1


1 2

2 $2t$

3 $2t^2$



DEFINITION--B-TREE HEIGHT

- $h \leq \log_t ((n+1)/2)$ ($t \geq 2$, page 489-490)
 - The worst case height is $O(\log n)$. Since the "branchiness" of a B-tree can be large compared to many other balanced tree structures, *the base of the logarithm* tends to be large.
 - Therefore, the number of nodes visited **during a search** tends to be smaller than the one required by other tree structures.
 - Although this does not affect the *asymptotic* worst case height, B-trees tend to have smaller heights than other trees with the same *asymptotic* height.
- 

OUTLINE

- Introduction
- Definition
- *Basic Operations*
- Applications



TWO CONVENTIONS

- The root of the B-tree is always in main memory, so that we never need to perform a DISK-READ on the root; we do have to perform a DISK-WRITE of the root, however, whenever the root node is changed.
- Any nodes that are passed as parameters must already have had a DISK-READ operation performed on them.



BASIC OPERATIONS

- *Searching a B-tree*
- Create an empty B-tree
- Splitting a node
- Inserting a key
- Deleting a key



BASIC OPERATIONS--SEARCH

- The search operation on a B-tree is *analogous to a search on a binary tree*.
- Instead of choosing between a left and a right child as in a binary tree, a B-tree search must make an $(n[x]+1)$ -way choice.
- The correct child is chosen by performing *a linear search* of the values in the node.



BASIC OPERATIONS--SEARCH


B-TREE-SEARCH (x, k)

```
1  $i \leftarrow 1$ 
2 while  $i \leq n[x]$  and  $k > \text{key}_i[x]$ 
3   do  $i \leftarrow i + 1$ 
4 if  $i \leq n[x]$  and  $k == \text{key}_i[x]$ 
5   then return ( $x, i$ )
6 elseif  $\text{leaf}[x]$ 
7   then return NIL
8 else Disk-Read( $c_i[x]$ )
9   return B-Tree-Search( $c_i[x], k$ )
```



BASIC OPERATIONS--SEARCH

B-Tree; **Minimum** Degree $t = 3$;
Minimum keys = 2; **Maximum** keys = 5



Search(21)



BASIC OPERATIONS—SEARCH ANALYSIS

- After finding the value greater than or equal to the desired value, the child pointer to the immediate *left* of that value is followed. If all values are less than the desired value, the *rightmost* child pointer is followed.
- Of course, the search can be terminated as soon as the desired node is found.
- $O(\log_t n)$ disk operation
- $O(t \log_t n)$ CPU time



BASIC OPERATIONS

- Searching a B-tree
- *Create an empty B-tree*
- Splitting a node
- Inserting a key
- Deleting a key



BASIC OPERATIONS--CREATE

- The *B-Tree-Create* operation creates an empty B-tree by allocating a new root node that has *no* keys and is a *leaf* node.
- Only the *root* node is permitted to have these properties; all other nodes must meet the criteria outlined previously.
- $O(1)$ disk operation
- $O(1)$ CPU time

```
B-Tree-Create(T)  
1  $x \leftarrow \text{Allocate-Node}()$   
2  $\text{leaf}[x] \leftarrow \text{TRUE}$   
3  $n[x] \leftarrow 0$   
4  $\text{Disk-Write}(x)$   
5  $\text{root}[T] \leftarrow x$ 
```



BASIC OPERATIONS

- Searching a B-tree
- Create an empty B-tree
- *Splitting a node*
- Inserting a key
- Deleting a key




BASIC OPERATIONS--WHY NEED SPLIT


- When inserting a key into a B-tree, we **can't** simply create a new leaf node and insert it, as the result tree would fail to be a **valid B-tree**.
- If a node becomes “too *full*” (having $2t-1$ keys), it is necessary to perform a split operation.
- It splits a full node y (*having $2t-1$ keys*) around its **median key** $key_t[y]$ into two nodes having $t-1$ keys each.
- The median key *moves up* into y 's **parent** to identify the dividing point between the two new trees.



BASIC OPERATIONS--SPLITTING



BASIC OPERATIONS--SPLITTING



BASIC OPERATIONS--SPLITTING

B-Tree-Split-Child(x, i, y)

```
1   $z \leftarrow \text{Allocate-Node}()$ 
2   $\text{leaf}[z] \leftarrow \text{leaf}[y]$ 
3   $n[z] \leftarrow t - 1$ 
4  for  $j \leftarrow 1$  to  $t - 1$ 
5      do  $\text{key}_j[z] \leftarrow \text{key}_{j+t}[y]$ 
6  if not  $\text{leaf}[y]$ 
7      then for  $j \leftarrow 1$  to  $t$ 
8          do  $c_j[z] \leftarrow c_{j+t}[y]$ 
9   $n[y] \leftarrow t - 1$ 
```

```
10 for  $j \leftarrow n[x] + 1$  downto  $i + 1$ 
11     do  $c_{j+1}[x] \leftarrow c_j[x]$ 
12  $c_{i+1}[x] \leftarrow z$ 
13 for  $j \leftarrow n[x]$  downto  $i$ 
14     do  $\text{key}_{j+1}[x] \leftarrow \text{key}_j[x]$ 
15  $\text{key}_i[x] \leftarrow \text{key}_t[y]$ 
16  $n[x] \leftarrow n[x] + 1$ 
17  $\text{Disk-Write}(y)$ 
18  $\text{Disk-Write}(z)$ 
19  $\text{Disk-Write}(x)$ 
```



BASIC OPERATIONS--SPLITTING ANALYSIS

- The split operation moves the median key of node y into its parent x where y is the i -th child of x .
- A new node, z , is allocated, and all keys in y right of the median key are moved to z . The keys left of the median key remain in the original node y .
- The new node, z , becomes the child immediately to the right of the median key that was moved to the parent x , and the original node, y , becomes the child immediately to the left of the median key that was moved into the parent x .

$c_i[x]$



BASIC OPERATIONS--SPLITTING ANALYSIS

- The split operation transforms a full node with $2t - 1$ keys into two nodes with $t - 1$ keys each. Note that one key is moved into the parent node.
- $O(1)$ disk operations
- $\Theta(t)$ CPU times



BASIC OPERATIONS

- Searching a B-tree
- Create an empty B-tree
- Splitting a node
- *Inserting a key*
- Deleting a key



BASIC OPERATIONS--INSERTION

- To perform an insertion on a B-tree, the appropriate node for the key must be located using an algorithm similar to *B-Tree-Search*.
- Next, the key must be inserted into the node.
 - If the node is **not full** prior to the insertion, no special action is required;
 - However, if the node is **full**, the node must be split to make room for the new key.
- Since **splitting** the node results in moving one key to the parent node, **the parent node** must not be full or another split operation is required.
- This process may repeat all the way up to the root and may require splitting the root node.

BASIC OPERATIONS--INSERTION


- This approach requires two passes:
 - The first pass locates the node where the key should be inserted (**Search**);
 - The second pass performs any required splits on the ancestor nodes (**Split**).
- We **don't wait** to find out whether we will need to split a full node. Instead, as we travel down the tree searching for the new key belongs, we split each full node we come to along the way.
- Thus, whenever we want to split a full node y , we are assumed that its parents is not full.



BASIC OPERATIONS--INSERTION

B-Tree-Insert(T, k)

```
1  $r \leftarrow \text{root}[T]$ 
2 if  $n[r] = 2t - 1$  then
3    $s \leftarrow \text{Allocate-Node}()$ 
4    $\text{root}[T] \leftarrow s$ 
5    $\text{leaf}[s] \leftarrow \text{FALSE}$ 
6    $n[s] \leftarrow 0$ 
7    $c_1 \leftarrow r$ 
8   B-Tree-Split-Child( $s, 1, r$ )
9   B-Tree-Insert-Nonfull( $s, k$ )
10 else
    B-Tree-Insert-Nonfull( $r, k$ )
```



BASIC OPERATIONS--INSERTION


B-Tree-Insert-Nonfull(x, k)

```
1   $i \leftarrow n[x]$ 
2  if leaf[ $x$ ]
3    then while  $i \geq 1$  and  $k < \text{key}_i[x]$ 
4        do  $\text{key}_{i+1}[x] \leftarrow \text{key}_i[x]$ 
5             $i \leftarrow i - 1$ 
6   $\text{key}_{i+1}[x] \leftarrow k$ 
7   $n[x] \leftarrow n[x] + 1$ 
8  Disk-Write( $x$ )
```


```
9 else while  $i \geq 1$  and  $k < \text{key}_i[x]$ 
10     do  $i \leftarrow i - 1$ 
11      $i \leftarrow i + 1$ 
12     Disk-Read( $c_i[x]$ )
13     if  $n[c_i[x]] == 2t - 1$ 
14         then
15             B-Tree-Split-Child( $x, i, c_i[x]$ )
16             if  $k > \text{key}_i[x]$ 
17                 then  $i \leftarrow i + 1$ 
18     B-Tree-Insert-Nonfull( $c_i[x], k$ )
```



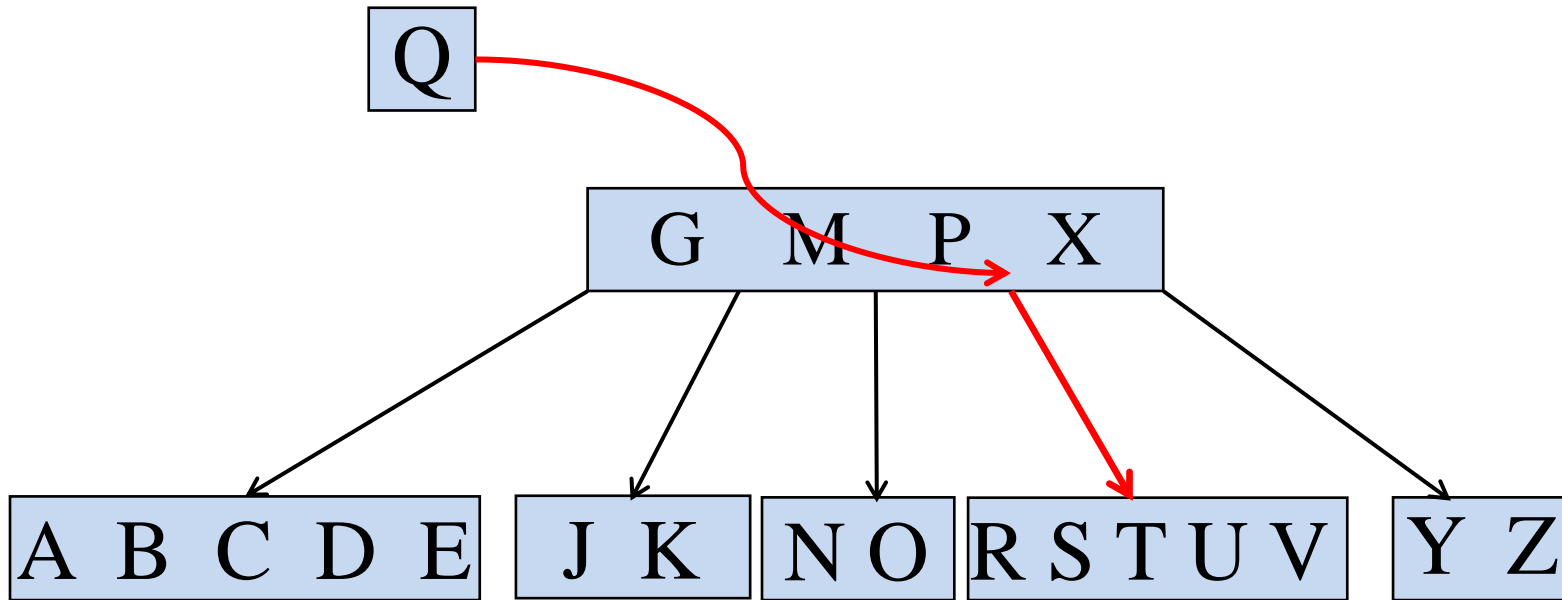
BASIC OPERATIONS--INSERTION




BASIC OPERATIONS--INSERTION



BASIC OPERATIONS--INSERTION



BASIC OPERATIONS--INSERTION

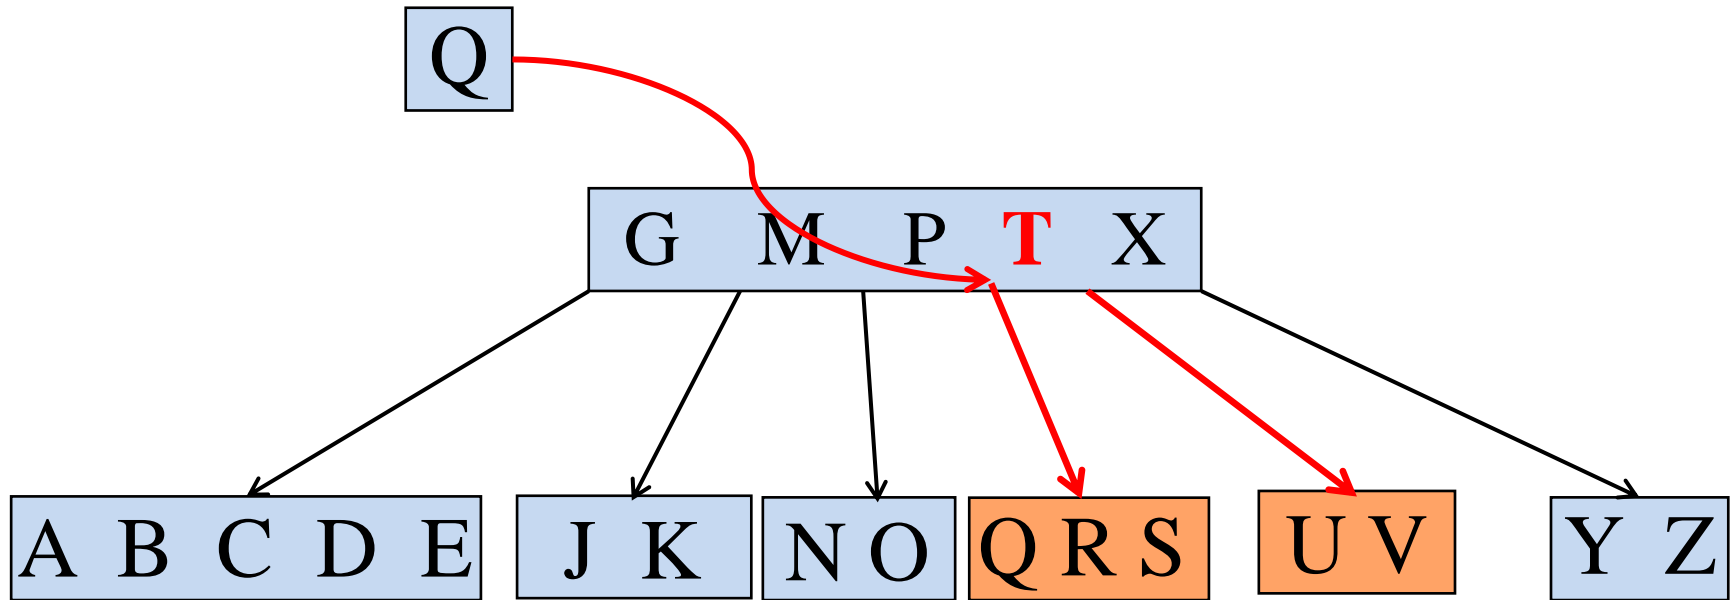


$$2t - 1$$

Minimum Degree $t = 3$



BASIC OPERATIONS--INSERTION

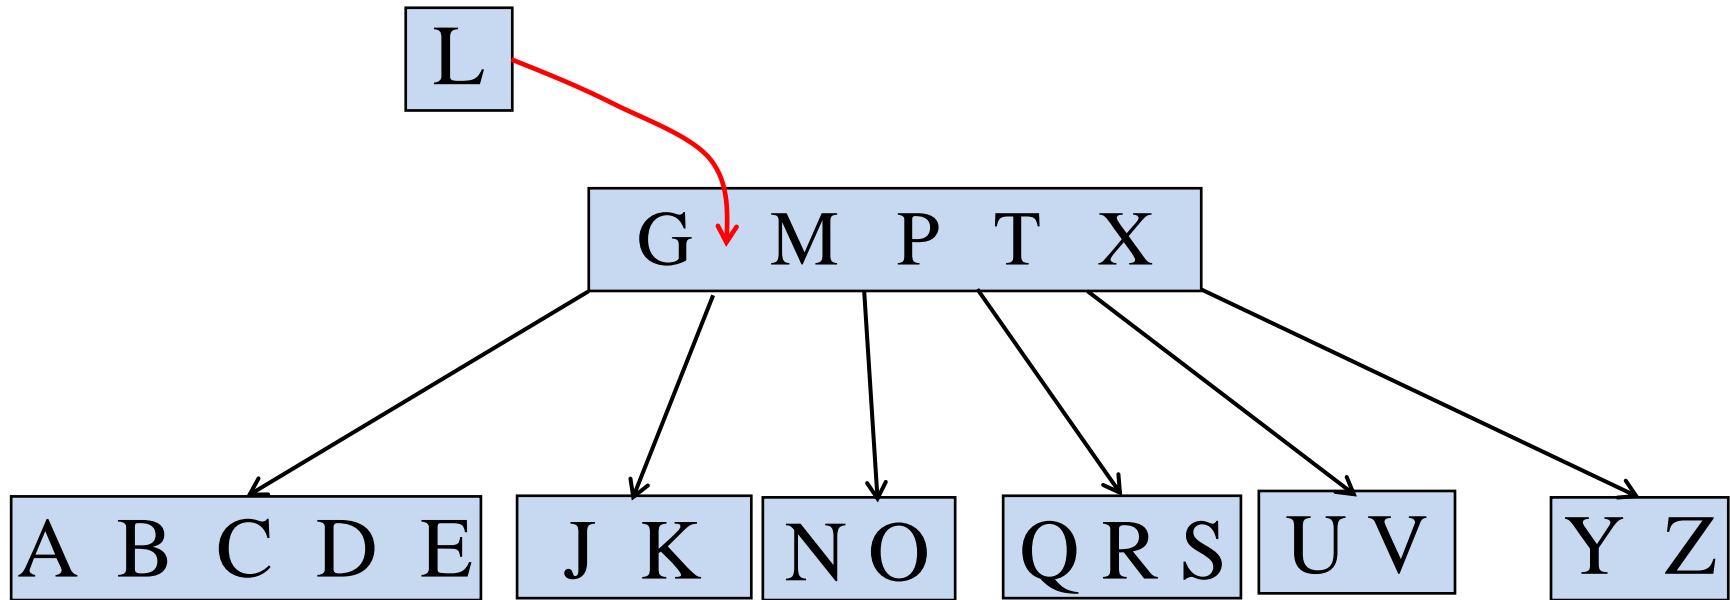


$$2t - 1$$

Minimum Degree $t = 3$



BASIC OPERATIONS--INSERTION




$$2t - 1$$

Minimum Degree $t = 3$



BASIC OPERATIONS--INSERTION




$$2t - 1$$


Minimum Degree $t = 3$



BASIC OPERATIONS--INSERTION



BASIC OPERATIONS--INSERTION




$$2t - 1$$

Minimum Degree $t = 3$



BASIC OPERATIONS--INSERTION




$$2t - 1$$

Minimum Degree $t = 3$



BASIC OPERATIONS--INSERTION



BASIC OPERATIONS--INSERTION

Sequence: (*a*, *g*, *f*, *b*, *k*, *d*, *h*, *m*, *j*, *e*, *s*, *i*, *r*, *x*, *c*, *l*, *n*, *t*, *u*, *p*) *t* = 5

(1)

<i>a</i>

(2)







<i>a</i>	<i>g</i>
----------	----------

(3)

<i>a</i>	<i>f</i>	<i>g</i>
----------	----------	----------


(4)

<i>a</i>	<i>b</i>	<i>f</i>	<i>g</i>
----------	----------	----------	----------




BASIC OPERATIONS--INSERTION


(10)




(11)






(12)



(13)




(14)




BASIC OPERATIONS--INSERTION


(15)




(16)




(17)



(18)






(19)



BASIC OPERATIONS--INSERTION

(20)



BASIC OPERATIONS--INSERTION ANALYSIS

- Since each **access to a node** may correspond to a costly disk access, it is desirable to avoid the second pass by ensuring that the parent node is never *full*.
- To accomplish this, the presented algorithm splits any full nodes encountered while descending the tree.
- Although this approach may result in unnecessary split operations, it guarantees that the parent never needs to be split and eliminates the need for a second pass up the tree.
- Since a split runs in linear time, it has little effect on the **$O(t \log_t n)$** running time of *B-Tree-Insert*.



BASIC OPERATIONS

- Searching a B-tree
- Create an empty B-tree
- Splitting a node
- Inserting a key
- *Deleting a key*



BASIC OPERATIONS--DELETION

- There are *two popular strategies* for deletion from a B-Tree.
 - Locate and delete the item, then restructure the tree to regain its invariants.
 - Do a single pass down the tree, but before entering (visiting) a node, restructure the tree so that once the key to be deleted is encountered, it can be deleted without triggering the need for any further restructuring.




BASIC OPERATIONS--DELETION

- **1.** If the key k is in node x and x is a leaf, delete the key k from x .
- **2.** If the key k is in node x and x is an internal node, do the following:
 - ***a.*** If the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the subtree rooted at y . Recursively delete k' , and replace k by k' in x . (Finding k' and deleting it can be performed in a single downward pass.)



BASIC OPERATIONS--DELETION

- **2.** If the key k is in node x and x is an internal node, do the following:
 - **b. (y has fewer than t keys)** Symmetrically, if the child z that follows k in node x has at least t keys, then find the successor k' of k in the subtree rooted at z . Recursively delete k , and replace k by k' in x . (Finding k' and deleting it can be performed in a single downward pass.)
 - **c.** Otherwise, if both y and z have only $t - 1$ keys, merge k and all of z into y , so that x loses both k and the pointer to z , and y now contains $2t - 1$ keys. Then, free z and recursively delete k from y .
- 

BASIC OPERATIONS--DELETION

- **3.** If the key k is not present in internal node x , determine the root $c_i[x]$ of the appropriate subtree that must contain k , if k is in the tree at all. If $c_i[x]$ has only $t - 1$ keys, execute step **3a** or **3b** as necessary to guarantee that we descend to a node containing at least t keys. Then, finish by recursing on the appropriate child of x .
 - **a.** If $c_i[x]$ has only $t - 1$ keys but has an immediate sibling with at least t keys, give $c_i[x]$ an extra key by moving a key from x down into $c_i[x]$, moving a key from $c_i[x]$'s immediate left or right sibling up into x , and moving the appropriate child pointer from the sibling into $c_i[x]$.



BASIC OPERATIONS--DELETION

- **3.** If the key k is not present in internal node x , determine the root $c_i[x]$ of the appropriate subtree that must contain k , if k is in the tree at all. If $c_i[x]$ has only $t - 1$ keys, execute step **3a** or **3b** as necessary to guarantee that we descend to a node containing at least t keys. Then, finish by recursing on the appropriate child of x .
 - **b.** If $c_i[x]$ and both of $c_i[x]$'s immediate siblings have $t - 1$ keys, merge $c_i[x]$ with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node.



BASIC OPERATIONS--VARIETIES

- **2-3-4 tree** is a B-tree of order 4
- **2-3 tree** is a B-tree that can contain only 2-nodes (a node with 1 field and 2 children) and 3-nodes (a node with 2 fields and 3 children)
- **B+ tree** (also known as a Quarternary Tree) is a variety of B-tree, in contrast to a B-tree, all records are stored at the lowest level of the tree; only keys are stored in interior blocks.
- **B*-tree** is a tree data structure, a variety of B-tree used in the HFS and Reiser4 file systems, which requires non-root nodes to be at least $2/3$ full instead of $1/2$.



OUTLINE

- Introduction
- Definition
- Basic Operations
- *Applications*




APPLICATIONS

- As databases cannot typically be maintained entirely in memory, B-trees are often used to index the data and to provide fast access.
 - For example, searching an unindexed and unsorted database containing n key values will have a worst case running time of $O(n)$.
- *Search Engine Index*
-



RUDOLF BAYER

- He is a professor emeritus of Informatics at the Technical University of Munich where he has been employed since 1972.
- He is noted for inventing three data sorting structures: the **B-tree** (with Edward M. McCreight), the **UB-tree** (with Volker Markl) and the **Red-black tree**.
- A recipient of 2001 ACM SIGMOD Edgar F. Codd Innovations Award



EDWARD M. MCCREIGHT

- Co-invented the **B-tree**, and improved Weiner's algorithm to compute the **suffix tree** of a string.
- Co-designed the *Xerox Alto workstation*.
- Co-led the design and construction of the Xerox Dorado computer at Xerox Palo Alto Research Center.
- Worked at Adobe Systems.
- Guest professor at the University of Washington, Stanford University, the Technical University of Munich, and the Swiss Federal Institute of Technology in Zürich.

